

Is Bravo actually faster to build in?

“Developing in Lora is hard, but developing in Bravo is faster and easier because the team now knows what to do.”

0 ISSUES ASSIGNED ON THE TWO BOARDS THE CLAIM CITES	1.83_x LORA’S TICKET VOLUME OVER BRAVO’S, JAN-AUG	2_d AUGUST MEDIAN CYCLE TIME – ON BOTH PLATFORMS	5→1 REPOSITORIES PER CHANGE – THE REAL EFFECT
---	---	--	---

Six claims, graded against the data

The DF4W and DF2W boards show Bravo development moving faster	● UNFALSIFIABLE	DF holds 24 issues, 23 of them Epics. D2W holds 22, all Epics. Every one unassigned, none resolved. These are product intake backlogs. Nothing on them can be timed.
Bravo development is faster	● NOT SUPPORTED	Bravo's real delivery project created ≈2,813 issues Jan–Aug against LORA's ≈5,144. August median created→resolved is 2 days on both . In June LORA was <i>faster</i> .
...because the team now knows what to do	● CONFOUNDED	Bravo's team has worked on this codebase since 2022 — 94 all-time authors, 23 active in 2026. That is tenure on a four-year-old codebase, not a platform being easier to learn.
Developing in Bravo is <i>easier</i>	● SUPPORTED	One change is fewer artefacts. A Bravo ticket decomposes into [BE] / [FE] / [QA] ; the LORA equivalent into five repositories and five deploys. Real, measurable, and the defensible core of the claim.
Bravo has a shallower onboarding curve	● HALF TRUE	Supported on skills — Java/Spring/Camunda are commodity. Unsupported on codebase: a new hire reads a 498k-LOC monolith with product identity in five places. The advantage is spent by week three.
DF2W is a live Bravo product delivering faster than LORA	● REFUTED	DF2W is fully configured in production and has started zero applications in 90 days. DF4W carries 5.5% of Bravo's volume. The two cited products are the smallest thing Bravo runs.

The claim is asking about ergonomics and being defended with velocity.

On velocity the data does not support it. On ergonomics it is right — and nobody had measured the size of the effect until now: five repositories per change against two lanes.

WHAT THE DATA REFUSES

Every throughput measure in this document favours LORA or ties. Bravo’s June median cycle time is **9 days against LORA’s 2**, and the cause is release batching, not architecture.

WHAT THE DATA CONFIRMS

One integration change is **1 repository and 1 deploy** in Bravo against **5 and 5** in LORA, in a mandated order. That is the whole real effect, and it is worth taking seriously.

The two boards the claim cites contain no development work

PROJECT	BOARD	ISSUES	TYPES	ASSIGNED	RESOLVED	FIRST → LAST CREATED
DF (DF4W)	2099, kanban	24	23 Epic, 1 Task	0	0	2025-12-12 → 2026-09-09
D2W (DF2W)	2877, kanban	22 of 36 keys	22 Epic	0	0	2026-06-02 → 2026-08-28
BL (LORA DP)	683	10,138 keys	Story, Sub-task, Task, Bug	assigned	thousands	2024-11-21 → 2026-09-10

Twelve of the 24 DF issues were created inside **four minutes** on 2025-12-12 — a backlog-loading session, not a sprint. D2W’s first sixteen were created inside **seventeen minutes**. Nine months after the first DF epic was written, not one has been assigned to anyone or moved out of To Do.

NOT A CRITICISM OF THE BOARDS

A product intake backlog is a legitimate artefact and both are well-formed ones. The point is what can be concluded from them: **a timeline view of these boards shows product intent, not delivery**. Any side-by-side putting board 2099 next to board 683 compares an epic backlog with a ticket stream.

Bravo’s development work lives in BLCS

The comparison the claim wants has to be made somewhere else. Recent BLCS titles confirm the scope: “#CapabilityTrue [4W Scoring] Adjustment Regular RAC 4W”, “[BE] Prepare Query Change GMB to NMH”, “[DF2W UW] Sync Surveyor Status Underwriting Return”.

BLCS – BRAVO LOS		BL – LORA DP
First issue	BLCS-985 , 2023-12-19	BL-3 , 2024-11-21
Latest at read time	BLCS-4811 , 2026-09-09	BL-10138 , 2026-09-10
Issue types in use	Epic, Story, Story Bug, Task, Sub-task, Bug	Epic, Story, Story Bug, Task, Sub-task, Bug
Sub-task convention	[BE] [FE] [QA] [DF2W UW]	[LSS] [LGS] [LPW] [LTW] [LTS] [LB0FE] [QA] [PR]


The taxonomies match, which is what makes the comparison possible — both projects run the same PMO template. Counting a BLCS ticket against a BL ticket is counting like against like at the *process* level. It is **not** like against like at the *change* level, which is where §6 matters.

DF2W WORK DOES APPEAR – IN BLCS

BLCS-4391 “[DF2W UW] Code Review - BM Return”, plus BLCS-4405 , 4420 , 4426 . So DF2W **is** being built. It is just not tracked on the DF2W board, and it has not yet taken a single production application.

Throughput: LORA carries 1.8× the ticket volume

BRAVO • BLCS LORA • BL



LORA’s lead narrows from 3.2× to 1.5× across the year. Jan–Aug totals: Bravo 2,813, LORA 5,144 — 1.83×.

METHOD

Jira issue keys are allocated monotonically per project, so the first key after a month boundary bounds the number created before it. Boundaries read directly: BLCS-1939 / BL-4795 at 1 Jan through BLCS-4752 / BL-9939 at 1 Sep.

CAVEAT, STATED PLAINLY

Key deltas are an **upper bound** — deleted and moved issues consume keys. D2W shows how real that is: **14 of 36 keys have no issue behind them**. The *trend* is robust to the error; the absolute ratio is not.

WHAT THIS DOES NOT MEAN

It does not mean LORA ships 1.8× the product change. Ticket count measures **decomposition** at least as much as output — which is precisely why the raw board comparison the claim invites is not a velocity measurement. See §6.

Cycle time: no durable advantage either way

Created → resolutiondate in calendar days, for every resolved issue created in a matched window. Two windows, four samples, n = 50 each, taken six months apart to avoid reading a single sprint.

SAMPLE	WINDOW	N	MEDIAN	MEAN	P90-ISH TAIL
Bravo BLCS	2026-06-01 → 06-08	50	9 d	22.3 d	60-98 d (9 issues)
LORA BL	2026-06-02 → 06-03	50	2 d	18.5 d	43-90 d (13 issues)
Bravo BLCS	2026-08-11 → 08-18	50	2 d	3.9 d	10-15 d (8 issues)
LORA BL	2026-08-11 → 08-12	50	2 d	5.3 d	14-29 d (9 issues)

1 • THE MEDIANS CONVERGE AT 2 DAYS

In August the two platforms are **indistinguishable** on the middle of the distribution. Whatever difficulty LORA imposes, it does not show up in how long a typical ticket takes to close.

2 • JUNE’S GAP RUNS THE WRONG WAY

Bravo 9 d against LORA 2 d — and the cause is visible: nine BLCS Stories created 2026-06-01 all resolved together on 07-17, 07-31 and 08-14. **Batch-closed on release.** That measures release cadence, not effort.

3 • SAME BIMODAL SHAPE ON BOTH

A mass of sub-3-day sub-tasks and a tail of Stories that close on release. That is **one delivery process running on two platforms** — exactly what §2 predicted.

Limits. Each sample covers one to six working days — a snapshot of two release cadences, not a distribution over the year. Windows differ in calendar length because BL creates ~1.8× the tickets, so 50 issues span fewer days. Both samples exclude unresolved issues, truncating the tail equally on both sides.

Team shape: 7 people against 22

DISTINCT ASSIGNEES, 100 SAMPLED ISSUES	BRAVO BLCS	LORA BL
Across both samples	7	22
June sample	6	19
August sample	7	13
Unassigned resolved issues	1	3

BRAVO’S STORY LAYER HAS ONE OWNER

In both samples essentially every Story is assigned to a **single account** while sub-tasks distribute across the remaining six — 9 of 9 long-lived Stories in June, 9 of 11 in August. Whether that account is a product owner, a lead or a de-facto integrator, the effect is the same: **the person who holds the shape of a Bravo change is a bus factor of one.** LORA’s Stories spread across a dozen accounts in the same samples.

7

ACTIVE ASSIGNEES ON A 498K-LOC MONOLITH CARRYING ~94% OF BFI ORIGINATION VOLUME

CROSS-CHECKED AGAINST THE CODEBASE

bravo-bpm-service has **94 all-time authors** but **23 active in 2026** across the whole service, of which the Scoring-and-Underwriting squad owns this repo. Seven active assignees is a thin shoulder — and it is the strongest argument in this document *against* moving new product families onto Bravo. Not because Bravo is slow, but because **there is nobody spare.**

HOW THIS BEARS ON THE CLAIM

“The team now knows what to do” is true, and it is true **of seven people.** LORA’s 22 assignees include people who joined this year and are shipping at the same median cycle time. That is evidence the LORA learning curve is *climbable* — not evidence it is absent.

The shape of one change — this is what “easier” actually means

This is the finding that survives every caveat, and the one the claim is really about. Take a single ordinary change: **bumping an upstream integration to v2**. In LORA it appears as six tickets across five repositories, all created within two minutes on 2026-08-11, all assigned to one engineer.

BL-9528	[LSS]	add one-obligor-summary-check-v2 schema
BL-9529	[LGS]	add one-obligor-summary-check-v2 handler
BL-9530	[LPW]	migrate get-obligor-summary to v2
BL-9531	[LTS]	register v2 proxy in authorization
BL-9532	[LBOFE]	migrate auto-exposure component to v2
BL-9533	[PR]	Review: one-obligor-summary-check-v2 refactor

That is the Schema-First order — LSS → LGS → LPW/LTW → LTS → LBOFE — mandated by LORA’s own service-dependencies rule, rendered as a ticket each. **It is not overhead someone added; it is the architecture made visible in the tracker.** The pattern repeats: BL-9506/9507/9508 split one fee calculation across [LSS] , [LTW] and [LBOFE] .

THE EQUIVALENT BRAVO CHANGE IS ONE SUB-TASK

BLCS-4433	[4W Scoring]	BE Implementation – Adjustment NST Rate – RO Rate Standard Override Indicator
-----------	--------------	---

OR, WHEN THE FRONT END IS INVOLVED, TWO LANES PLUS QA

BLCS-4388	[FE]	Development – DF2W section & field visibility on Ringkasan Data
BLCS-4389	[FE]	Manual Testing – verify DF2W reduced view + 4W no-regression
BLCS-4390	[FE]	Deploy – Deploy to SIT

One repository and one deploy, against five and five

PER ONE INTEGRATION CHANGE	BRAVO	LORA
Repositories touched	1 – <code>bravo-bpm-service</code>	5 – LSS, LGS, LPW, LTS, LBOFE
Deploys to land it	1	up to 5, in a mandated order
Jira artefacts	1-3	5-6
Backward-compatibility obligation	none in-process	additive-only – 8 worker versions run concurrently
Cost when the change is wrong	one revert	a version-ordered unwind

This is the entire defensible content of “faster and easier”, and it is worth taking seriously.

A monolith with a relational aggregate really is fewer moving parts per change than thirteen repositories with a schema-first ordering rule. It also explains §3 without any appeal to productivity: LORA’s ticket count is ~1.8× Bravo’s and its per-change ticket count is ~2-5× Bravo’s — so **per change**, the two teams may well be shipping at similar rates with very different bookkeeping.

WHAT THE FIVE-REPO TAX BUYS

Product isolation at the data layer, schema-validated contracts, and the ability to add an automated step with **no orchestration edit at all** — 172 activities against 3 hard precursors.

WHAT THE ONE-REPO CHANGE PAYS

One job executor, one deployment, and a BPMN edit that **ships for every product at once**.

Neither is free. The claim is right that Bravo’s per-change cost is lower. It is silent on what that cost buys.

The onboarding curve

	BRAVO	LORA
Hiring pool	Java 17 / Spring Boot / Camunda 7 – commodity	Go + a bespoke GSM planner + Temporal – “you can get Java or Golang developers, but they won’t be able to immediately read the code”
Written onboarding	none found in this analysis pack	docs/onboarding/ – 14 files, ~5,550 lines, with working hands-on tooling
Week-1 curriculum	not established	also not established – LORA’s own people.md grades this “overstated, but the entry point is inverted”
What a hire must actually read	498k LOC main Java; product identity in 5 places; 197 setStatus sites in 73 files; a 10,419-line surveyor assignment service	~617k LOC Go across 13 repos; 172 activity constructors; 155 preconditions with no generated index
Concepts with no analogue elsewhere	BPMN escalation as a return channel; a per-application jsonb activity on/off matrix; five-place product discrimination	ReadSet/WriteSet planning; determinism constraints; schema-versioned task queues

The commodity-skills advantage is real and it is **front-loaded**. It gets a hire to their first compile faster. It does not help with the parts of Bravo that are actually hard, and by the evidence of this pack those are large: an aggregate root with no behaviour, a lifecycle smeared across four status vocabularies, and orchestration untested by 1,453 of 1,457 test files.

Neither platform has a week-1 curriculum. That is the finding both teams share — and it is cheaper to fix than either architecture.

What the claim gets right, and what it does not

“Developing in Lora is hard”	● PARTLY TRUE	And <i>specifically</i> true: five repositories, a mandated ordering rule, additive-only schema changes, eight concurrent worker versions. Measured in §6.
“Developing in Bravo is faster”	● NOT SUPPORTED	Fewer tickets (1.83×), identical August median cycle time, slower June median. No throughput measure in this document favours Bravo.
“...and easier”	● SUPPORTED	One repository, one deploy, no cross-service ordering. This is the real effect and it is worth about 3–5 artefacts per change.
“...because the team now knows what to do”	● CONFOUNDED BY TENURE	Bravo’s seven active assignees have been on this codebase for up to four years. LORA’s twenty-two include this year’s joiners closing tickets at the same median. Knowing what to do is a property of the people, and the seven do not scale.

A FIFTH THING THE CLAIM DOES NOT SAY, AND SHOULD

The two products it cites are the two **smallest** things Bravo runs. DF4W is the only product on the unified spine — **5.5%** of started applications over 90 days. DF2W is fully configured and has taken **zero**. If those are the reference implementations for “Bravo is easier”, then the easiness has not yet been tested at volume, on the monoliths, where **~94% of the book actually runs** and where 2026 saw more commits than the spine did.

Six actions

4 Name Bravo’s bus factor · S, URGENT Seven active assignees, with the Story layer held by one account, on the service carrying ~94% of origination. Before any new product family is placed on Bravo, that number has to be a staffing decision someone has made on purpose.	1 Stop citing boards 2099 and 2877 as delivery evidence · S Zero assigned, zero resolved. Either wire DF/D2W epics to their BLCS children so the timeline reflects delivery, or state explicitly that Bravo delivery is tracked in BLCS and compare there.
2 Measure changes, not tickets · S–M The 1.83× ticket ratio and the 5-repo decomposition point the same direction and cancel. Agree one unit — merged PRs per product change, or Story-level lead time — and publish it for both. Until then neither team’s velocity claim is checkable.	3 Fix the LORA per-change tax directly · M It is the true part of the complaint. BL-9528. .9532 are mechanical and ordered; a generator that scaffolds schema, handler, activity migration, authorization entry and FE call from one spec turns six tickets into one plus review. LORA’s own people.md designs this as “Tier 0 — a script not a prompt”. It is not built.
5 Write the week-1 curriculum for both · S EACH Neither platform has one. LORA has the raw material and the wrong entry point; Bravo has commodity skills and no map of a 498k-LOC monolith.	6 Re-ask the question when DF2W has volume · M DF2W is the cleanest available test of “building a new product on Bravo is easy”: fully configured, zero applications, 22 unstarted epics. Its actual cost — first BLCS ticket to first production application — is the number this debate needs and does not have.

30, 60, 90 days

Almost none of this is Bravo engineering time — it is engineering management, PMO and, for one item, the LORA squad. That matters, because \$5 is the constraint on everything else in this pack. This plan spends about **half a day** of Squad S&U’s time.

SEQUENCING RULE

The **staffing decision comes first**, because it sets the capacity every other document’s plan assumes.

Days 0–30 SETTLE WHAT IS MEASURED, AND WHO DOES IT 4 Name the bus factor. Decide the staffing and the standing allocation for remediation work. <i>Leadership · a meeting, day 1.</i> Done when a recorded decision exists, including what percentage of Squad S&U time is available for the other four plans. 1 Board hygiene. Wire DF/D2W epics to their BLCS children, or say so in the board description. <i>PMO · 1 d.</i> 3 Spike the LORA scaffolding generator. <i>LORA squad · 5 d spike.</i> A prototype generates the BL-9528. .9532 shape.	Days 31–60 AGREE THE UNIT OF COMPARISON 2 Measure changes, not tickets. Agree one unit and publish it for both platforms. <i>EM + LORA squad · 3 d.</i> Done when “Bravo shipped 40 product changes last month, LORA 44” is a sentence someone can check — and the 1.83× ticket ratio stops being quoted as velocity.	Days 61–90 REMOVE THE TAX, FIX ONBOARDING 3 Ship the scaffolding generator. <i>LORA squad · 15 d.</i> The strongest argument for Bravo is answered on the merits, not rebutted. 5 Week-1 curriculum, Bravo — a map of the monolith and its five product-discrimination mechanisms. <i>EM + S&U · 3 d.</i> 5 Week-1 curriculum, LORA — invert the entry point from the SDK to a loan moving. <i>LORA EM · 3 d.</i> 6 Start the DF2W clock. Instrument first- BLCS - ticket to first-production-application. <i>PMO · 1 d.</i>	Deferred WITH A TRIGGER 6 Re-ask the “is Bravo easier” question with real DF2W data. DF2W has zero applications today and 22 unstarted epics — there is nothing to measure yet. <i>Trigger: when DF2W takes production volume, on the clock started in phase 3.</i>
--	---	---	--

What changes when this is done

RECOMMENDATION	EXAMPLE WHEN DONE
Delivery tracked where delivery happens	A DF4W epic on board 2099 shows its <code>BLCS</code> children and a real start date. The timeline stops reading as nine months of inactivity.
One agreed unit of change	“Bravo shipped 40 product changes last month, LORA shipped 44” is a sentence someone can check, instead of two ticket counts that measure different things.
Scaffolding for the five-repo change	<code>BL-9528..9532</code> becomes one ticket and a generated PR set. The strongest argument for Bravo loses most of its force, on the merits.
Bus factor named	Placing a new product family on Bravo is a decision that includes “and we will hire three more engineers” – or it is a decision not taken.
Week-1 curriculum	A Go hire reaches a moving loan in week 1 instead of the SDK; a Java hire reaches a moving loan instead of <code>SurveyorAssignmentServiceImpl</code> .
DF2W measured end to end	The claim gets a real number attached to it, on the product it was made about.

On velocity, the claim is not supported. On ergonomics, it is right — and the fix is a generator, not a migration.